


If we run this clustering algorithm using labeled data, we can assess the performance of our label propagation

Suppose we have 20 images:

- 4 - Ryan
- 5 - Melanie
- 6 - Katie
- 5 - Johnny

Our algorithm should predict 2 clusters.

We can also measure some pairwise metrics to see how we are doing.



Most of the Ryan-nodes were given the same label (3). This is what we want!

Unfortunately, one Ryan-node was given a label-2, which is associated with a Melanie cluster.

How can we detect these errors?

We will use pairwise metrics

• tally if: labels match + truth match: **LT**

labels match + truth disagree $L \ T$
labels disagree + truth match $L \ T$
for a pair of nodes

- Loop through all pairs of nodes in your graph and continue tallying these up

$$\text{pairwise precision} : \frac{(L \ T)}{(L \ T) + (L \ T)}$$

high precision: conservative clustering
low precision: aggressive clustering

$$\text{pairwise recall} : \frac{(L \ T)}{(L \ T) + (L \ T)}$$

high recall: inclusive clusters
low recall: dispersive clusters